

ECE492/592 – Operating Systems Design: Project #4

Due date: December 10, 2018

Note: You can perform this assignment in teams of two.

Objectives

- To understand Xinu's memory management.
- To understand how to implement a virtual memory system. This includes understanding in details interactions between the hardware and the software.
- To understand how interrupt handling works (including the related hardware and software interactions).

Overview

The goal of this project is to implement virtual memory management and demand paging in Xinu. Refer to the chapters of the Intel Manual (available in Moodle) related to Paging and Interrupt Handling. In particular, as you read the Intel Manual, focus on the following: (i) registers used to support virtual memory and paging, (ii) PDE and PTE format, (iii) interrupt handling support.

Specifically, you will find the information you need in the **Intel Architecture Software Developer's Manual Volume 3: System Programming Guide**:

- Memory management:
 - o Memory management registers – Section 2.4
 - o Invalidating TLBs – Section 2.6.4
 - o Paging and Virtual Memory – Section 3.6
 - o Translation Lookaside Buffers – Section 3.7
- Interrupt handling: Chapter 5

In addition, see attached slides **xinu_memory_mng.pdf** for an overview of Xinu's memory management.

Please read carefully the whole project description below before starting.

System call implementation

Your solution must implement the system calls listed below.

```
1. pid32 vcreate (void *funcaddr, uint32 ssize, uint32 hsize, pri16 priority,
                char *name, int nargs, ...)
```

This system call creates a process with a specified heap size. The process's heap must be private and exist in the process's own virtual memory space. Parameter `hsize` indicates the heap size (in number of pages). **Use a 4KB page size.** The *total* number of pages available for heap is determined by macro `MAX_HEAP_SIZE` defined in `paging.h`. Note that `MAX_HEAP_SIZE` is not a per-process limitation – it indicates the total amount of heap space available to all processes. `MAX_HEAP_SIZE` does not include the memory allocated through the `getmem` system call.

Note: While processes created with Xinu's `create` function should not have a private virtual heap, enabling virtual memory and paging might require you to modify the `create` function as well.

2. `char* vmalloc (uint32 nbytes)`

This function allocates the desired amount of memory (in bytes) off a process's virtual heap space, and returns `SYSERR` if the allocation fails.

3. `syscall vfree (char* ptr, uint32 nbytes)`

This function frees heap space (previously allocated with `vmalloc`). `vfree` returns `OK` in case of success, and `SYSERR` in case of failure. In case of failure, none of the pages involved should be freed.

Additional requirements

1. Memory initialization

The original Xinu memory initialization has a bug. In addition, it does not allow the whole 4GB address space to be available for paging. Replace the original `meminit.c` and `i386.c` files in the `system` folder with the ones attached, which fix the problem. You are recommended to look into these files and understand what has been fixed. As you can see in Xinu's code, besides the TEXT, DATA, and BSS regions, Xinu statically reserves a so-called HOLE area for boot-loading purposes.

Important hint: After paging is enabled, all memory accesses are through virtual addresses. Therefore, you need to map the static segments (TEXT, DATA, etc.) in the virtual memory space of each process.

2. Free Frame Space (FFS)

The FFS is the physical memory space where processes map their virtual heap space. FFS must be released when no longer required (in other words, heap frames should be released upon heap deallocation). The total amount of FFS frames available is determined by macro `MAX_FSS_SIZE` defined in `paging.h`. `MAX_FSS_SIZE` is not a per-process limitation – it indicates the total amount of FSS available to all processes. *When a user process maps a FFS frame to a virtual page, that FFS frame should not be visible to other user processes* (i.e., user processes can map to their virtual address space only the FFS frames they use).

3. Disk Space Simulation [only students taking course at graduate level are required to implement swapping]

Virtual memory typically uses disk space to extend the physical memory of the machine. However, the Virtual Box version of Xinu that you are using does not have a file system. Thus, you need to simulate the disk space using memory. Reserve a total of `MAX_SWAP_SIZE` frames to emulate the disk space. Macro `MAX_SWAP_SIZE` is defined in `paging.h`. *This disk space must not be mapped onto the virtual memory space of the user processes.*

4. Page directory and Page Tables

Page directories and page tables must be always *resident* in physical memory (i.e., they should never be swapped to disk). Physical memory used by page directories and page tables must be released when no longer required. A total of `MAX_PT_SIZE` frames can be reserved for page directory and page tables.

[Additional requirement for students taking course at the graduate level] - *The area of memory where page directories and page tables are stored cannot be accessed by user processes (i.e., this area must not be mapped to user processes' virtual address space).*

5. Heap allocation

You must use a **lazy allocation policy** for heap allocation. That is, the physical space should be reserved not at allocation time, but when the virtual page is first accessed. Accesses to pages that have not been previously allocated must cause a `SEGMENTATION_FAULT`.

6. Page replacement and swapping

You must use the **random replacement policy**. In addition, you must have **global replacement** (that is, a process might cause the eviction of pages belonging to other processes). On page eviction, only dirty pages must be copied to disk.

Simplifying assumptions

1. You can assume that the operating system performs allocations and triggers faults **at a page granularity**.

For example, assuming 8-byte pages, the following allocation operations:

```
char *ptr1 = vmalloc(16);  
char *ptr2 = vmalloc(15);
```

will be treated the same way. In other words, the OS will not keep track of the number of bytes within a page that have been effectively allocated. In addition, the OS won't trigger a segmentation fault if the code accesses `ptr2[15]` even if the allocation was just for 15 bytes.

2. Similarly, you can assume that allocation operations **are aligned to the beginning of a page**. For example, you can assume that the following allocations:

```
char *ptr1 = vmalloc(16);  
char *ptr2 = vmalloc(15);
```

will go to two different pages.

Hints

1. Paging should be enabled at the end of system initialization. A good place to enable it is the end of the `sysinit()` function in `initialize.c`
2. x86 PDE and PTE have bits that are reserved for software use. You can use those bits as you please in your implementation.
3. **Interrupt handling** – Study the code in `clkinit.c`, `clkdisp.S` and `clkhandler.c` to find out how to install an interrupt service routine (ISR) in Xinu. Implement the page fault handler in the following two files: `pagefault_handler_disp.S` and `pagefault_handler.c`. As indicated in the Intel Manual, a page fault triggers interrupt 14. Thus, you need to install your page fault handler to interrupt 14. Intel

processors automatically push an error code on the stack when an ISR is called, so you need to pop out the error code off the stack within your Page Fault Handler.

4. For convenience, we have added a `control_reg.c` and `paging.h` that you can extend.
5. Use the same page table for all system processes.
6. Even if Xinu does not have a clear difference between user mode and system mode, you should aim at a design that conceptually distinguished “user mode” and “kernel mode”. To this end, assume that, when a user process issues a system call, that call is executed in kernel mode. On the other hand, the rest of the user process’s code (that is, the code within the function executed by the process) is executed in user mode.
7. **Test your code “incrementally”.**
 - A. Start by testing the code with paging enabled but without invoking the three system calls above. This will allow verifying that the page directories and page tables have been setup correctly. If this is not the case, the system will hang on startup. Add a function to dump the content of page directories and page tables in a format that is easy for you to read and interpret.
 - B. When (A) works, test user process creation and termination. Start with a single process.
 - C. Add calls to `vmalloc` and `vfree` from one process. Do not introduce swapping.
 - D. Add heap accesses (to trigger lazy allocation).
 - E. Add more processes, so to test context switch between user processes.
 - F. Add swapping.

We attach a simple test case file (`testcases.c`). You can use it and modify it to implement test cases following (A-F).

Requirement for extra credits (10 points)

Provide an implementation where user processes have a virtual stack, too. The maximum number of stack frames that can be allocated should be defined in macro `MAX_STACK_SIZE` that you can add to `paging.h`.

Note: it is **not** allowed to post questions on the extra credits problems on the class mailing list. Work on these problems only if you can do it individually or within your team.

Submissions instructions

1. **Design Document:** You need to submit a document covering the various aspects of your design **by November 27**. This document should cover only the required portion of the project. If you decide to implement the requirement for extra credits, describe the related aspects of your design in the final report only. We attach a draft indicating the aspects that your document should cover. Submit your design document in Moodle by November 27.
2. **Report:** Your report should:
 - describe any changes in design over your original design document (no need to document the aspects you had already discussed in your design document);
 - explain what works and what not in your implementation;
 - indicate if your implementation also covers the extra points requirement (in this case, indicate the basic aspects of the design).
3. **Test cases:** For this project, you do **not** need to submit your test cases. We will test your implementations using different test cases (i.e., different `main.c` files). Therefore, do not implement any essential

functionality (other than your test cases) in the `main.c` file. Also, turn off debugging output before submitting your code. Our test cases will be similar to (and an extension of) the ones in `testcases.c`.

4. **Extra points:** Submit **only one implementation**. You can choose to either submit a code that implements only the basic requirements (for no extra points), or a code that also implements the requirements for extra points.
 - Submit the implementation for extra points only if you have tested it extensively and are certain that it works properly. In other words, make sure that all the working features of your basic implementation still operate correctly (otherwise you will lose credits).
 - If only one team member worked on the extra points, make two separate submissions (two codes and reports): one team member will submit the implementation of the required portion, and the other team member will submit the implementation covering also the requirements for extra points.
5. Go to the `xinu/compile` directory and invoke `make clean`.
6. As for previous project, create a `xinu/tmp` folder and **copy** all the files you have modified/created into it (the `tmp` folder should have the same directory structure as the `xinu` folder).
7. Go to the parent folder of the `xinu` folder. Compress the whole `xinu` directory into a `tgz` file.

```
tar czf xinu_project4.tgz xinu
```

8. Submit report and code through Moodle by December 10.